

# AC-Key: Adaptive Caching for LSM-based Key-Value Stores

Fenggang Wu, Ming-Hong Yang, Baoquan Zhang, and David H.C. Du

ATC '20

**Presenter: Kai Li**

# Background

- In enterprise workloads, read operations exhibit “**hot spots**” in LSM-tree-based KVSs (LSM-KVS) for both point lookups and range queries.
- **Caching** “hot spots” can improve read performance.
- Three type of entries that can be cached in LSM-KVS: Block, Key-Value (KV), and Key-Pointer (KP).

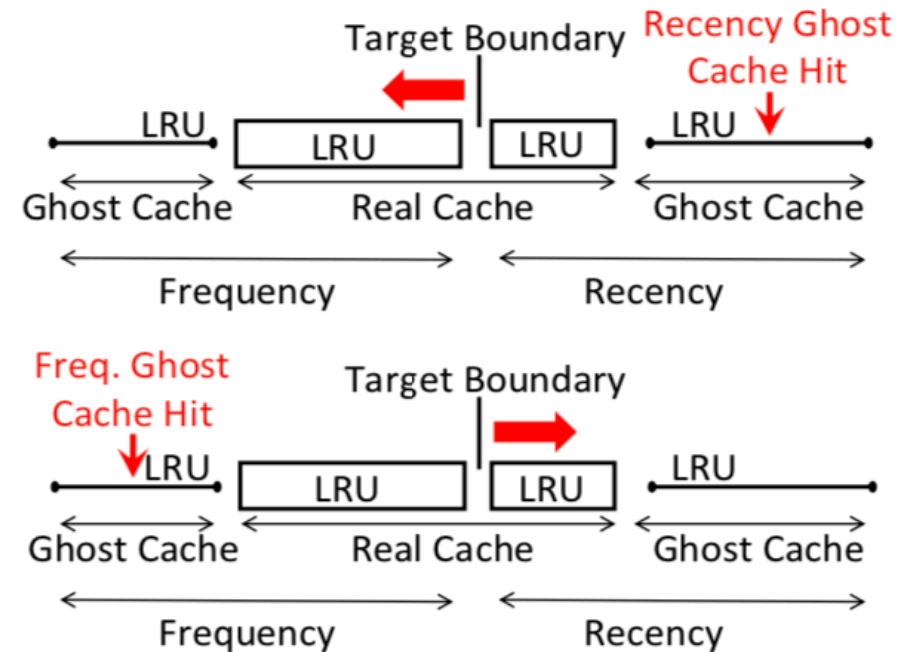
# Motivation

- Caching LSM-KVS is **challenging**
  - Data have different sizes/levels -> different cache costs/benefits
    - Cache a record at a deeper level can bring more benefits, as it reduces more storage IOs
  - Different type of read: point lookup, range query
    - **Cache a block** favors range query
    - **Cache a KV/KP** favors point query
    - When value size is **large**
      - Cache a KP is more space-efficient than cache a KV

	Block	KV	KP	Point	Range	Adaptive
LevelDB	Yes	No	No	Inefficient	Supported	Fix-sized
RocksDB	Yes	Yes	No	Large Value inefficient	Supported	Fix-sized
Cassandra	No	Yes	Yes	Efficient	Not Supported	Fix-sized
AC-Key	Yes	Yes	Yes	Efficient	Supported	Adaptive-sized

# Cache Size Adjustment with Ghost Cache

- Adaptive Replacement Cache (ARC)
  - Real Cache + Ghost Cache
  - Ghost Cache hit will push the boundary
  - Adaptive to different cache components




**Figure 3:** ARC Algorithm.

# AC-Key: Hierarchical Adaptive Caching

Upper Level: Point Cache vs Block Cache

Lower level: KV Cache vs KP cache



# Get Handling

- *Case I:* Hit in KV Cache
  - The value is returned without any I/O incurred.
- *Case II:* Miss in KV Cache but hit in KP Cache.
  - Check Block Cache (and bring to Block Cache if missing)
  - Promote KP Cache to KV Cache
- *Case III:* Miss both in KV Cache and KP Cache.
  - Search every sorted run level by level
  - Move into KP Cache

# Flush and Compaction

- Sync the caches only in flushing time, not during *Put*
- AC-Key updates KP and block Caches when compaction affects any of cached KP entries or blocks

# Caching Efficiency Factor

- To quantitatively analyze the trade-off between the costs and benefits of the cache entries
  - Target Boundary Adjustment by *Caching Efficiency Factor*  $E$
  - Adjustment  $\Delta = kE$

$$E = \frac{b}{s}, \quad (1)$$

where:  $E$  = caching efficiency factor of one cached entry,  
 $b$  = number of saved storage I/O if cached,  
 $s$  = caching space taken by this entry.

$$b = \begin{cases} 1 & \text{if block,} \\ f(m) & \text{if KV entry,} \\ f(m) - 1 & \text{if KP entry.} \end{cases} \quad (2)$$

where:  $m$  = number of SSTs to search for the key,  
 $f(m)$  = number I/Os to get a key. It is a function of  $m$ .


$$m = \begin{cases} n_0/2 & \text{if } l = 0, \\ l + n_0 & \text{if } l \geq 1. \end{cases} \quad (3)$$

where:  $n_0$  = max number of SSTs  $L_0$  can hold,  
 $l$  = the level where the key resides.

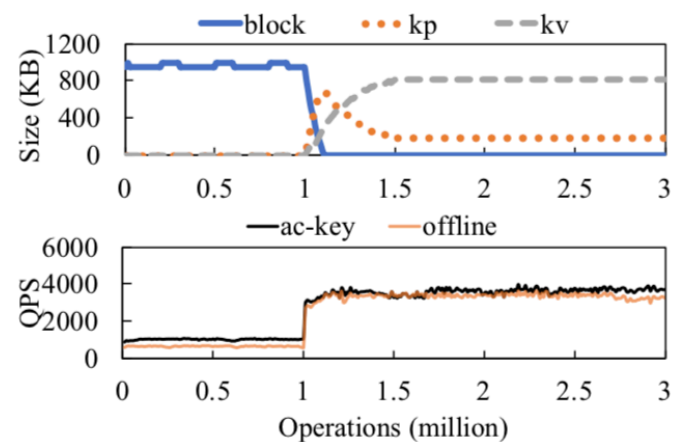

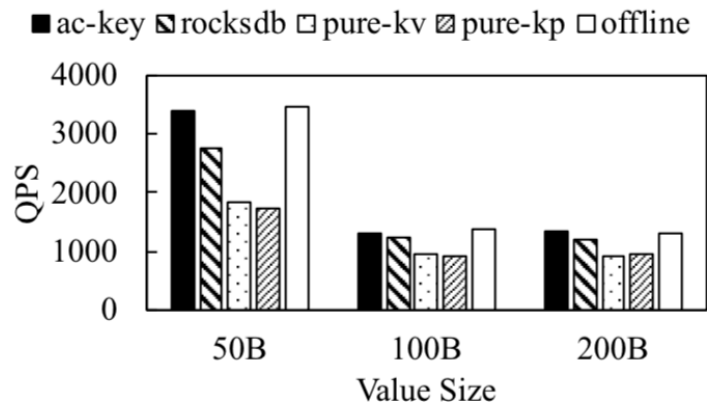




# Evaluation

- Implementation on RocksDB, 5.6K LoC.
- Compare *AC-Key* with *pure-kv*, *pure-kp*, *rocksdb*, *offline*



# Extensive Evaluations





Thank you!

Q&A